

HW1_dbagchi2

Diptendra Nath Bagchi (dbagchi2@illinois.edu)

03 February, 2020

Question 1


a) Multivariate Gaussian Simulation

```
set.seed(1)
epsilon = rnorm(n = 1000, mean = 0, sd = 1)
mvr_generation = function(mu, sigma) {
  no_of_variates = length(mu)
  lower_tri = t(chol(sigma))
  std_random = rnorm(no_of_variates)
  (lower_tri %*% std_random + mu)
}
mu = c(0, 1, 2)
sigma = rbind(c(1.0, 0.5, 0.5),
              c(0.5, 1.0, 0.5),
              c(0.5, 0.5, 1.0))
multi_matrix = matrix(0, nrow = 1000, ncol = 3)
for (i in 1:1000) {
  multi_matrix[i, ] = mvr_generation(mu = mu, sigma = sigma)
}

# y is generated as per
y = 0.25 * multi_matrix[, 1] + 0.5 * multi_matrix[, 2] + multi_matrix[, 3] + epsilon

plot(x = multi_matrix[, 3], y = y,
     main = "Dependent vs independent",
     xlab = "Independent Value (X3)",
     ylab = "Dependent value (Y)",
     col = "deepskyblue")
```

Dependent vs independent



Proof:

$$\begin{aligned} Cov(X) &= Cov(\mu + CZ) = Cov(CZ) = C.cov(Z).C^t = CIC^t = \Sigma \\ mean(X) &= mean(\mu + CZ) = mean(\mu) + mean(CZ) = \mu + 0 = \mu \end{aligned}$$

b) Function to implement own K-Nearest Neighbours

```
myknn = function(xtest, xtrain, ytrain, k){
  n_test = nrow(xtest)
  n_pred = numeric(n_test)
  iter = seq(1, n_test)
  for(i in iter){
    xtest_i = xtest[i, ]
    distance <- rowSums(abs(sweep(xtrain, 2, xtest_i)))
    dist_y = cbind(distance, ytrain)
    min_dist = dist_y[order(dist_y[, 1], decreasing=FALSE),]
    n_pred[i]=sum(min_dist[1:k, 2]) / k
  }
  n_pred
}
```

c) Calculating the mse from KNN for k = 5

```
# Dividing the data into training and testing data
xtrain = multi_matrix[1:400, ]
ytrain = y[1:400]
xtest = multi_matrix[401:nrow(multi_matrix), ]
```

```
ytest = y[401:length(y)]
knn_pred = myknn(xtest, xtrain, ytrain, 5)
```

```
calc_mse = function(predicted, actual) {
  mean((actual - predicted) ^ 2)
}
```

```
calc_mse(knn_pred, ytest)
```

```
## [1] 1.342291
```

The mean squared error is 1.34.

d) What is the optimal tuning parameter?

```
k = c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 200, 300, 400)
```

```
knn_test_pred = lapply(k, myknn, xtest = xtest, xtrain = xtrain, ytrain = ytrain)
```


```
mse_test = sapply(knn_test_pred, calc_mse, actual = ytest)
```

```
knn_train_pred = lapply(k, myknn, xtest = xtrain, xtrain = xtrain, ytrain = ytrain)
```

```
mse_train = sapply(knn_train_pred, calc_mse, actual = ytrain)
```

```
dof = 400 / k
plot(dof, mse_train, type = "b", col = "red",
     ylim = c(0, 5), main = "KNN Train/Test MSE",
     xlab = "Degrees of Freedom (N/K)",
     ylab = "Mean Squared Error"
    )
lines(dof, mse_test, col = "deepskyblue", type = "b")
legend("topright", legend=c("Training Error", "Testing Error"),
     col=c("red", "deepskyblue"), lty=1:1, cex=0.8)
```

KNN Train/Test MSE



The optimal tuning parameter is `r_k[which.min(mse_test)]`.

e) Linear Regression

```
train_df = data.frame("x" = xtrain, "y" = ytrain)
lmod = lm(y ~ x.1 + x.2 + x.3, data = train_df)
lmod_pred = predict(object = lmod, data.frame("x" = xtest))
lmod_mse = calc_mse(predicted = lmod_pred, ytest)
```

Degree of freedom of linear model is p which is much less than the dof of KNN $400/9 \sim 45$ - which makes KNN a complex model than the linear regression. Also, as the dependent variable is defined as a linear combination of the independent variables - which also makes the linear model a better choice. The mse of linear regression is 1.1688589.

f) Try a new model

```
y = 0.25 * multi_matrix[, 1] + 0.5 * multi_matrix[, 2] + (multi_matrix[, 3] ^ 2) + epsilon

xtrain = multi_matrix[1:400, ]
ytrain = y[1:400]
xtest = multi_matrix[401:nrow(multi_matrix), ]
ytest = y[401:length(y)]

k = c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 200, 300, 400)

knn_test_pred = lapply(k, myknn, xtest = xtest, xtrain = xtrain, ytrain = ytrain)

mse_test = sapply(knn_test_pred, calc_mse, actual = ytest)
```

```

knn_train_pred = lapply(k, myknn, xtest = xtrain, xtrain = xtrain, ytrain = ytrain)

mse_train = sapply(knn_train_pred, calc_mse, actual = ytrain)

train_df = data.frame("x" = xtrain, "y" = ytrain)
lmod = lm(y ~ x.1 + x.2 + x.3, data = train_df)
lmod_pred = predict(object = lmod, data.frame("x" = xtest))
lmod_mse = calc_mse(predicted = lmod_pred, ytest)

```

KNN model is performing better than K-Nearest Neighbours by a significant margin. The test mse of knn is 2.377 and the test mse of linear model is 3.5763353. Also the specification of the model is incorrect as there is square terms which are missing from the linear model but in KNN it does not effect much because of its non-parametric nature. Hence, in this case KNN performs better than the linear model.

Question 2

```

# part 1
set.seed(1)
x = seq(from = 1, to = 97, by = 1)
add_97 = lapply(x, function(x) {rnorm(1000, mean = 0, sd = 1)})
add_97_type_cast = matrix(unlist(add_97), ncol = 97)

x_matrix_a = cbind(multi_matrix, add_97_type_cast)
y2_a = y

# part 2
set.seed(1)
matrix_a_temp = lapply(x, function(x) {runif(3, min = 0, max = 1)})
matrix_a_temp = matrix(unlist(matrix_a_temp), ncol = 97)
x_trans_a = multi_matrix %*% matrix_a_temp
x_matrix_b = cbind(multi_matrix, x_trans_a)
y2_b = y

myknn = function(xtest, xtrain, ytrain, k) {
  n_test = nrow(xtest)
  n_pred = numeric(n_test)
  iter = seq(1, n_test)
  for(i in iter){
    xtest_i = xtest[i, ]
    distance <- rowSums(abs(sweep(xtrain, 2, xtest_i)) ^ 2)
    dist_y = cbind(distance, ytrain)
    min_dist = dist_y[order(dist_y[, 1], decreasing=FALSE),]
    n_pred[i]=sum(min_dist[1:k, 2]) / k
  }
  n_pred
}

# running knn for both the covariates
xtrain = x_matrix_a[1:400, ]
ytrain = y2_a[1:400]
xtest = x_matrix_a[401:nrow(x_matrix_a), ]
ytest = y2_a[401:length(y2_a)]

k = c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 200, 300, 400)

```

```

knn_test_pred = lapply(k, myknn, xtest = xtest, xtrain = xtrain, ytrain = ytrain)

mse_test = sapply(knn_test_pred, calc_mse, actual = ytest)

knn_train_pred = lapply(k, myknn, xtest = xtrain, xtrain = xtrain, ytrain = ytrain)

mse_train = sapply(knn_train_pred, calc_mse, actual = ytrain)

```

The optimal K for part I is 10 and the test MSE is 15.6818354.

```

xtrain = x_matrix_b[1:400, ]
ytrain = y2_b[1:400]
xtest = x_matrix_b[401:nrow(x_matrix_b), ]
ytest = y2_b[401:length(y2_b)]

k = c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 200, 300, 400)

knn_test_pred = lapply(k, myknn, xtest = xtest, xtrain = xtrain, ytrain = ytrain)

mse_test = sapply(knn_test_pred, calc_mse, actual = ytest)

knn_train_pred = lapply(k, myknn, xtest = xtrain, xtrain = xtrain, ytrain = ytrain)

mse_train = sapply(knn_train_pred, calc_mse, actual = ytrain)

```

The optimal K for part II is 2 and the test MSE is 2.7566101.

Part II outperforms part I primarily because of the fact that the covariates are highly correlated in with the original three variables and hence, the mse is lower than the other part because in that case all the covariates are independent and therefore the original x's are unable to explain the model fully.

Question 3

```

library(ElemStatLearn)
trn = as.data.frame(zip.train)
tst = as.data.frame(zip.test)

df = rbind(trn, tst)
df$V1 = as.factor(as.character(df$V1))

nfold = 10
index_fold = sample(rep(1:nfold, length.out = nrow(df)))

all_k = c(seq(from = 10, to = 100, by = 10))

accu_matrix = matrix(NA, length(all_k), nfold)

for (l in 1:nfold) {
  for (k in 1:length(all_k)) {
    knn_fit = kknn(V1 ~ ., train = df[index_fold != l, ],
                   test = df[index_fold == l, ], k = all_k[k])
    accu_matrix[k, l] = mean(knn_fit$fitted.values == df$V1[index_fold == l])
  }
}

```

```

    }
  }

  best_k = all_k[which.min(apply(accur_matrix, 1, mean))]

  find_closest = function(xtest, xtrain, ytrain, k) {
    ytest = rep(0, nrow(xtest))
    ids = rep(0, nrow(xtrain))
    dist = rep(0, nrow(xtrain))
    for (i in 1:nrow(xtrain)) {
      ids[i] = i
      dist[i] = sum((xtest[1, 2:257] - xtrain[i, 2:257]) ^ 2)
    }
    df = data.frame("ids" = ids, "distance" = dist)
    df = df[order(df$distance), ][1:k+1, ]
    return(df)
  }

  temp = find_closest(df[1, ], df, df[, 1], best_k)

  ids = temp$ids

  par(mfrow=c(10,10), mar=c(1,1,1,1))
  for(i in 1:best_k) {
    image(zip2image(rbind(zip.train, zip.test), temp$id[i]), col=gray(256:0/256), zlim=c(0,1), xlab="", y
  }


## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"

```

[illegible]

[illegible]

```
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
## [1] "digit 6 taken"
```

1357	855	5335	988	4095	7300	8790	2783	6017	3773
									
8005	2583	2506	6431	5596	6266	3737	4292	5600	5827
									
3084	3757	2705	847	7590	8291	484	7174	2786	9128
									
3712	4717	6665	3227	518	973	807	651	1058	856
									
5265	2737	6632	8064	1040	8139	4459	8939	3008	2078
									
8000	8744	9114	118	7570	6308	3485	3037	6959	6961
									
2443	8532	1593	2457	428	5056	5257	6630	4228	7625
									
4677	4943	8112	7636	3034	5033	288	7316	337	2380
									
414	4559	1128	6317	7589	4742	5503	6113	2542	6108
									
4407	6470	7982	2422	4607	5057	1558	8802	6190	3566
									

The observation is correctly classified using K-Nearest Neighbours based on the optimal value of k.

KNN seems to perform well on this data set as the representation of the points (images of hand written) are similar in the data set. It means 6 is very similarly written in all the images of 6 and hence the pixel values are also very similarly spread out in the images. Therefore, the distances are close to other numbers and hence KNN performs well on this data set.